



LibSci for Accelerators (libsci_acc)

- **Provide basic libraries for accelerators, tuned for Cray**
- **Must be independent to OpenACC, but fully compatible**
- **Multiple use case support**
 - Get the base use of accelerators with no code change
 - Get extreme performance of GPU with or without code change
 - Extra tools for support of complex code
- **Incorporate the existing GPU libraries into libsci**
- **Provide additional performance and usability**
- **Maintain the Standard APIs where possible!**



Why libsci_acc?

- **Code modification is required to use existing GPU libraries!**
- **Several scientific library packages are already there**
 - CUBLAS, CUFFT, CUSPARSE (NVIDIA), MAGMA (U Tennessee), CULA (EM Photonics)
- **No Compatibility to Legacy APIs**
 - cublasDgemm(...)
 - magma_dgetrf(...)
 - culaDgetrf(...)
 - Why not dgemm(), dgetrf()?
- **Not focused on Fortran API (C/C++)**
 - Require CUDA data types, primitives and functions in order to call them
- **Performance**

Three Interfaces For Three Use Cases

- Simple interface

`dgetrf(M, N, A, lda, ipiv, &info)`

`dgetrf(M, N, d_A, lda, ipiv, &info)`

GPU + CPU

CPU

GPU

- Device interface

`dgetrf_acc(M, N, d_A, lda, ipiv, &info)`

- CPU interface

`dgetrf_cpu(M, N, A, lda, ipiv, &info)`

Simple Interface

- **You can pass either host pointers or device pointers to the simple interface**
- **Host memory pointer**
 - Performs hybrid operation on GPU
 - If problem is too small, performs host operation
- **Device memory pointer**
 - Performs operation on GPU
- **BLAS 1 and 2 perform computation local to the data location**
 - CPU-GPU data transfer is too expensive to exploit hybrid execution

Device Interface

- Device interface gives higher degree of control
- Requires that you have already copied your data to the device memory
- API
 - Every routine in libsci has a version with `_acc` suffix
 - E.g. `dgetrf_acc`
 - This resembles standard API except for the suffix and the device pointers

CPU Interface

- **Sometimes apps may want to force ops on the CPU**
 - Need to preserve GPU memory
 - Want to perform something in parallel
 - Don't want to incur transfer cost for a small op
- **Can force any operation to occur on CPU with `_cpu` version**
- **Every routine has a `_cpu` entry-point**
- **API is exactly standard otherwise**

libsci_acc Usage - Basics

- Supports Cray and GNU compilers.
- Fortran and C interfaces (column-major assumed)
 - Load the module *craype-accel-nvidia35*.
 - Compile as normal (dynamic libraries used)
- To enable threading in the CPU library, set **OMP_NUM_THREADS**
 - E.g. export OMP_NUM_THREADS=16
- Assign 1 single MPI process per node
 - Multiple processes cannot share the single GPU
- Execute your code as normal

libsci_acc DGEMM Example

- Starting with a code that relies on dgemm.
- The library will check the parameters at runtime.
- If the size of the matrix multiply is large enough, the library will run it on the GPU, handling all data movement behind the scenes.
- **NOTE:** Input and Output data are in CPU memory.

```
call dgemm('n','n',m,n,k,alpha,&  
          a,lda,b,ldb,beta,c,ldc)
```


libsci_acc Interaction with OpenACC

- If the rest of the code uses OpenACC, it's possible to use the library with directives.
- All data management performed by OpenACC.
- Calls the device version of dgemm.
- All data is in CPU memory before and after data region.

```
!$acc data copy(a,b,c)
```

```
!$acc parallel
```

```
!Do Something
```

```
!$acc end parallel
```

```
!$acc host_data use_device(a,b,c)
```

```
call dgemm_acc('n','n',m,n,k,&
               alpha,a,lda,&
               b,ldb,beta,c,ldc)
```

```
!$acc end host_data
```

```
!$acc end data
```

libsci_acc Interaction with OpenACC

- libsci_acc is a bit smarter than this.
- Since 'a,' 'b', and 'c' are device arrays, the library knows it should run on the device.
- So just dgemm is sufficient.

```
!$acc data copy(a,b,c)
```

```
!$acc parallel
```

```
!Do Something
```

```
!$acc end parallel
```

```
!$acc host_data use_device(a,b,c)
```

```
call dgemm      ('n','n',m,n,k,&
                  alpha,a,lda,&
                  b,ldb,beta,c,ldc)
```

```
!$acc end host_data
```

```
!$acc end data
```

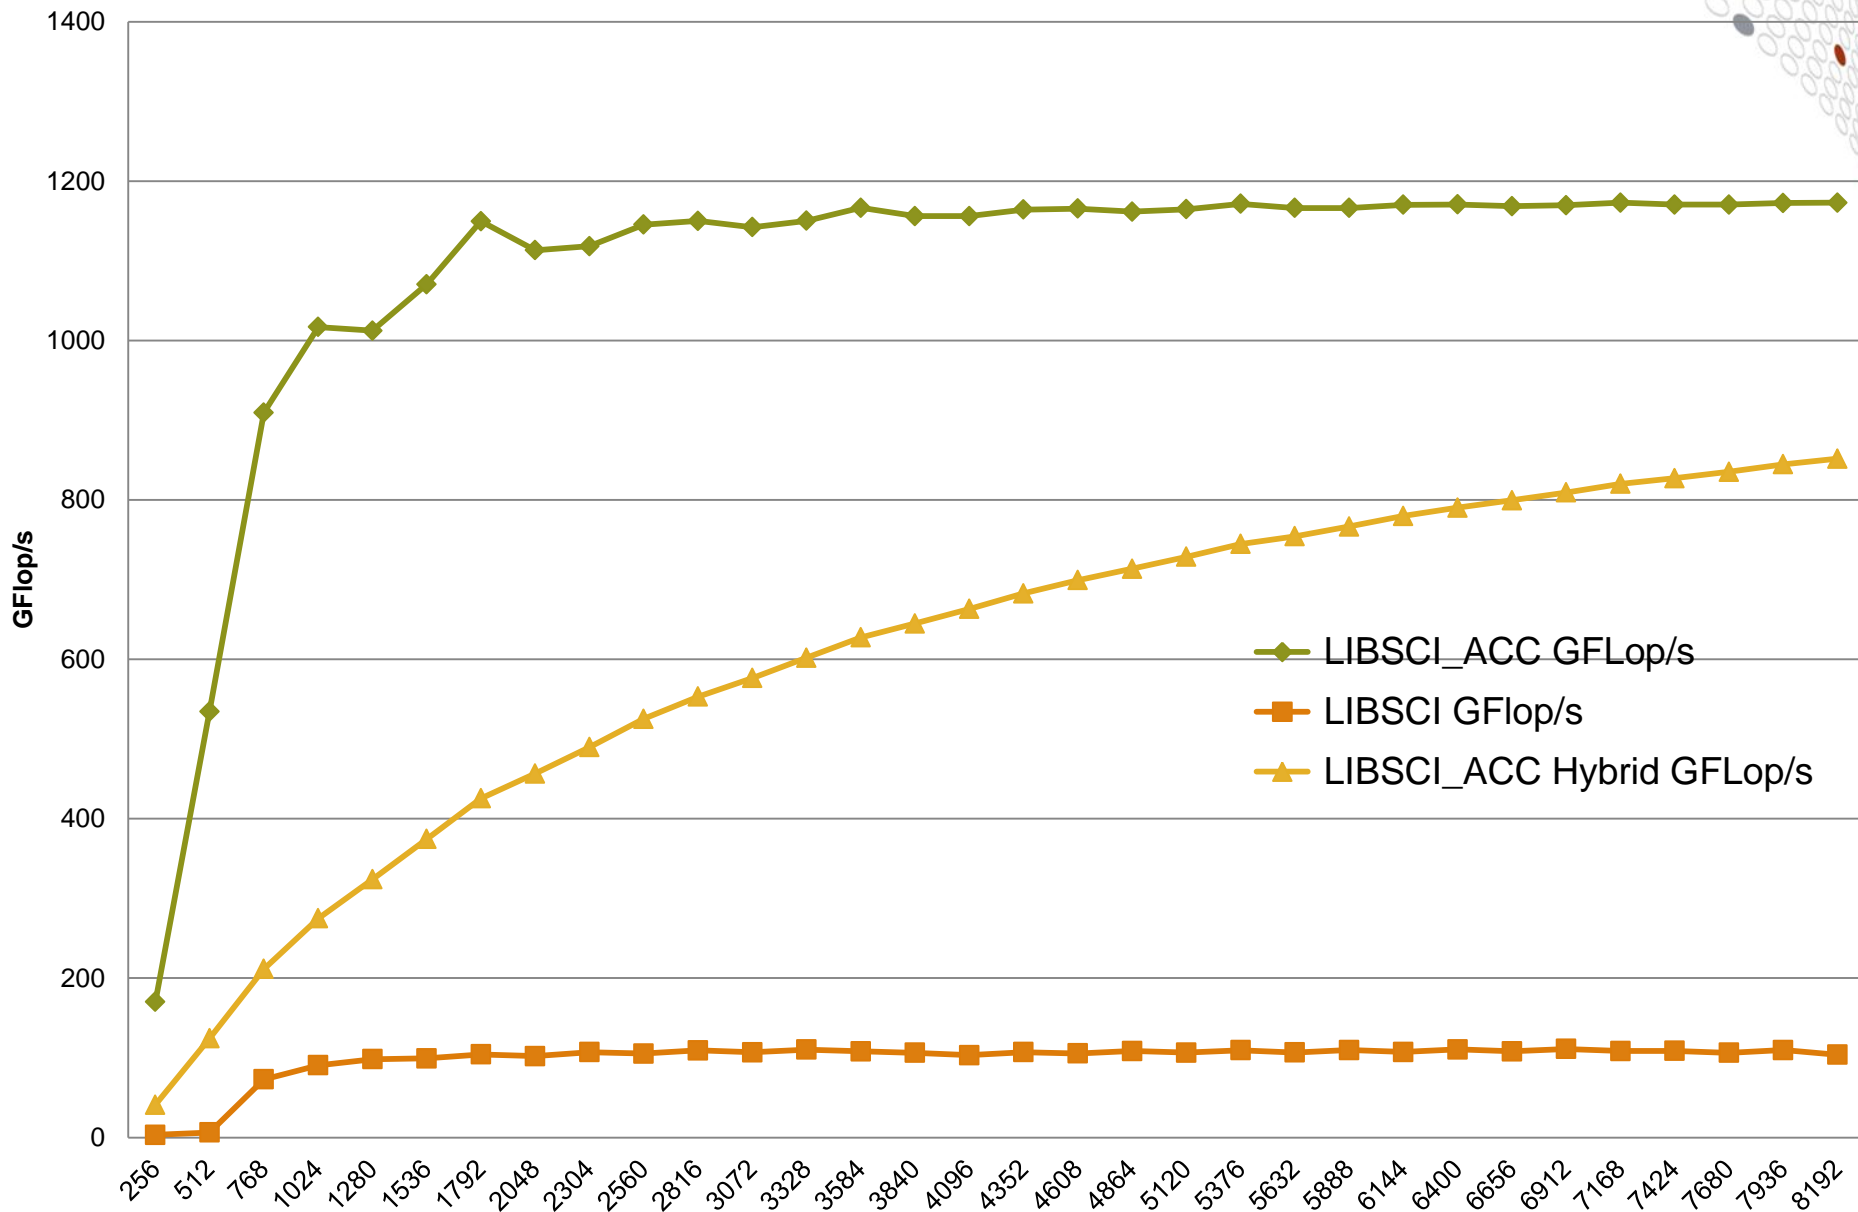
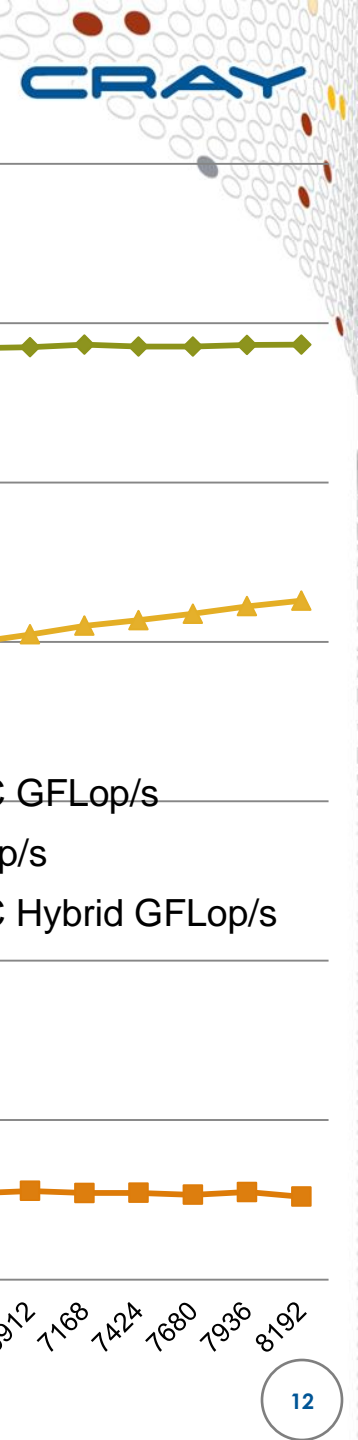


Advanced Controls

- The communication avoidance (CA) version of DGETRF/ZGETRF can be enabled by setting the environment variable `LIBSCI_ACC_DLU = CALU / LIBSCI_ACC_ZLU = CALU`
- **Change Split Ratio of Hybrid GEMM routines**
 - `LIBSCI_SGEMM_SPLIT=0.9`
 - `LIBSCI_DGEMM_SPLIT=0.8`
 - `LIBSCI_CGEMM_SPLIT=0.9`
 - `LIBSCI_ZGEMM_SPLIT=0.8`
- **Force simple API to always call CPU routine**
 - `CRAY_LIBSCI_ACC_MODE=2`

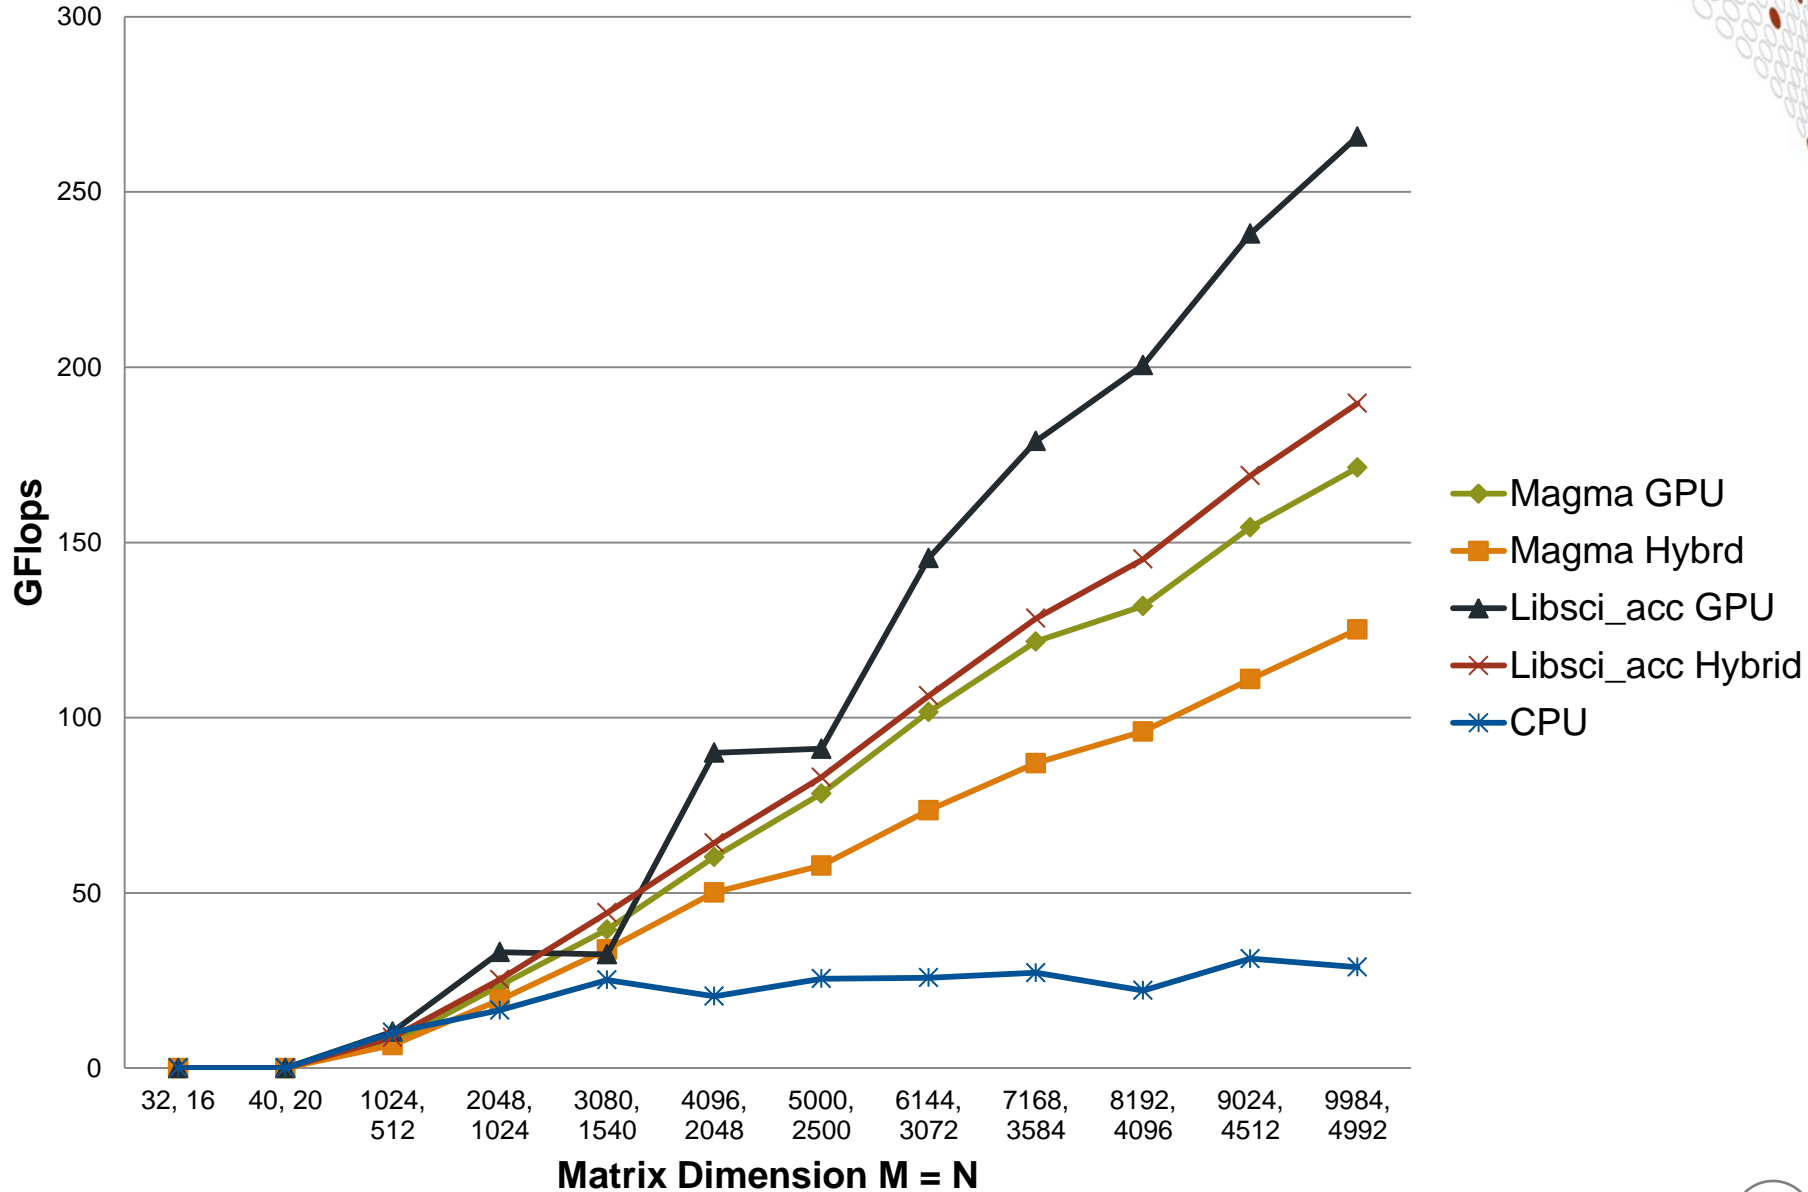
Matrix Multiplication :: Double (DGEMM)

XK7 Kepler :: Nov 2012



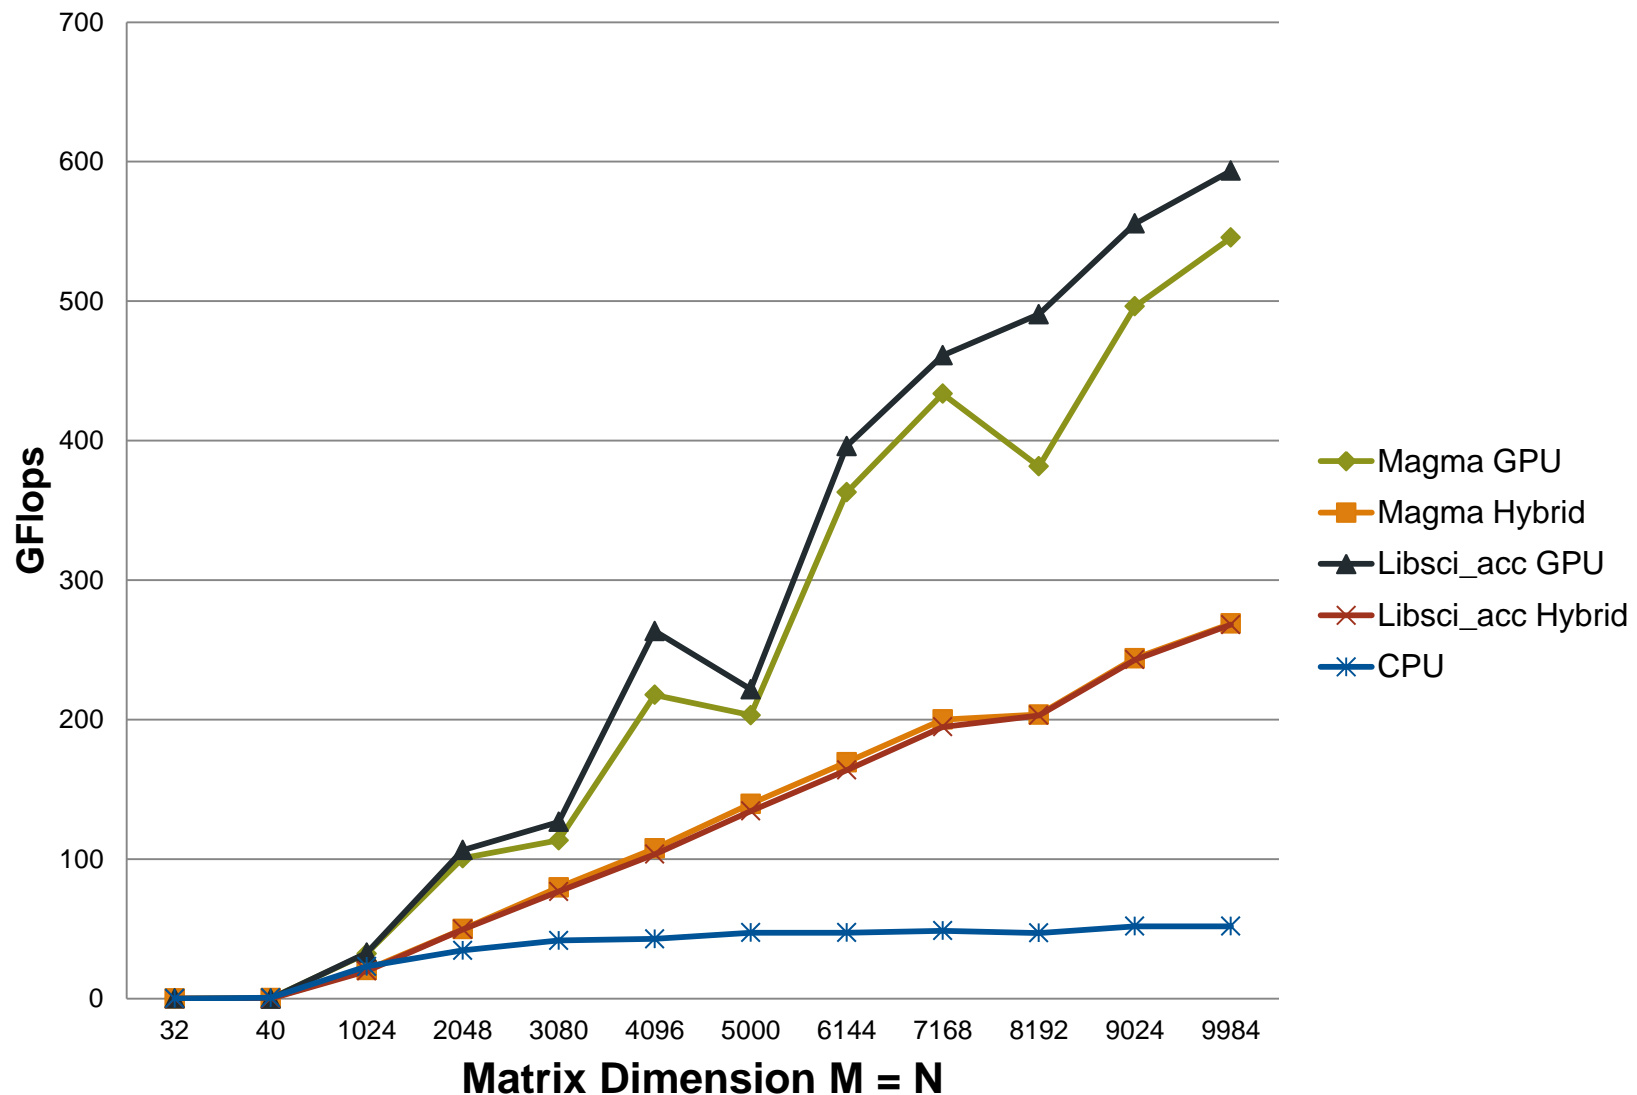
LAPACK QR factorization :: DGEQRF

XK7 Kepler :: Nov 2012



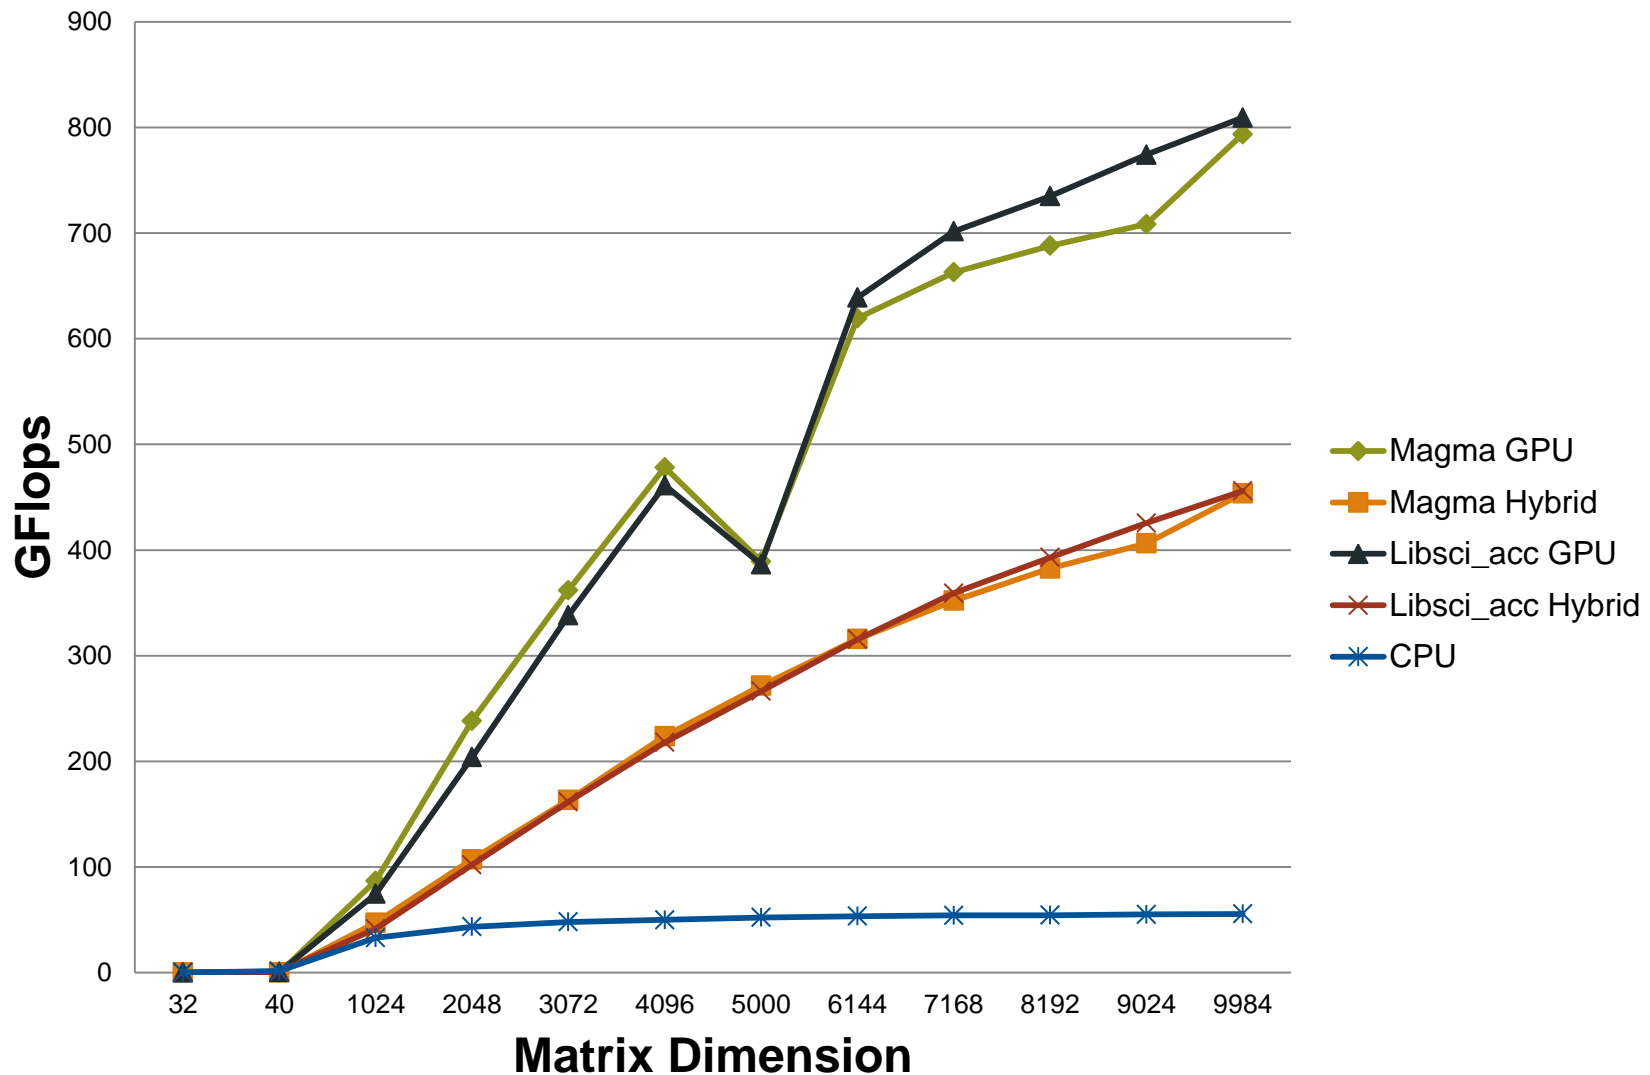
LAPACK LU factorization :: double (DGETRF)

XK7 Kepler :: Nov 2012



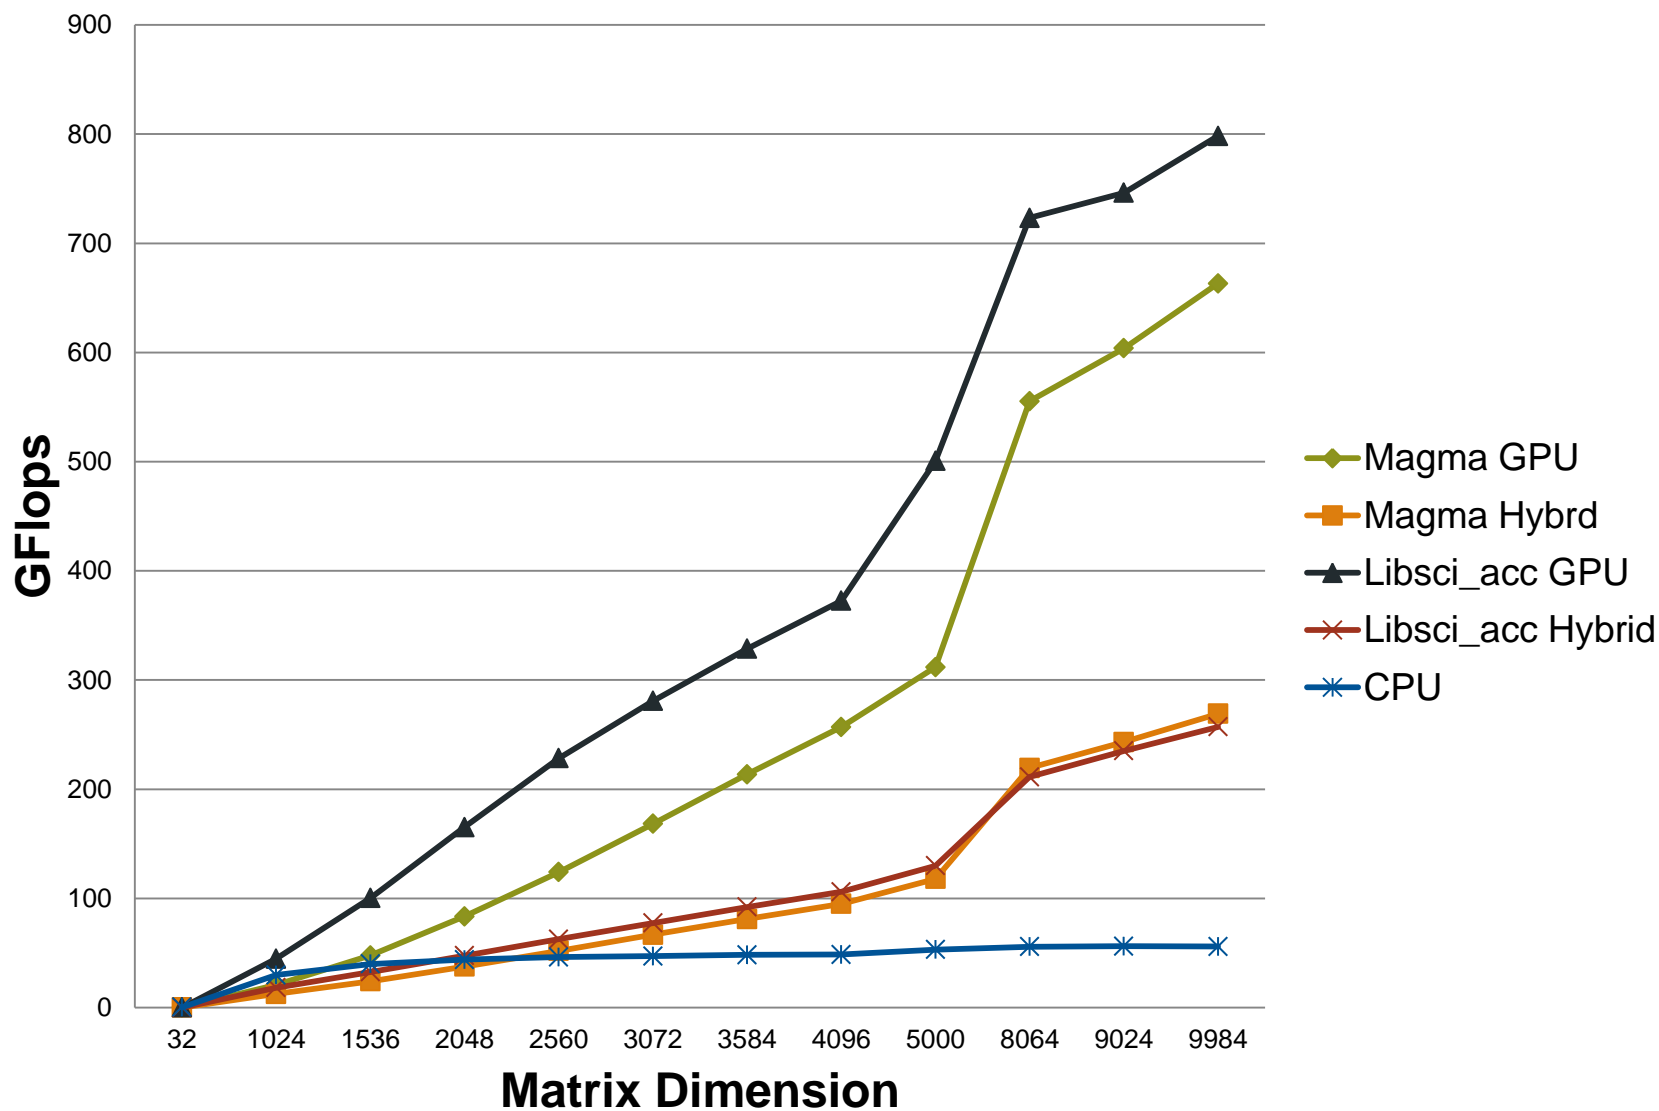
LAPACK LU factorization :: double complex (ZGETRF)

XK7 Kepler :: Nov 2012



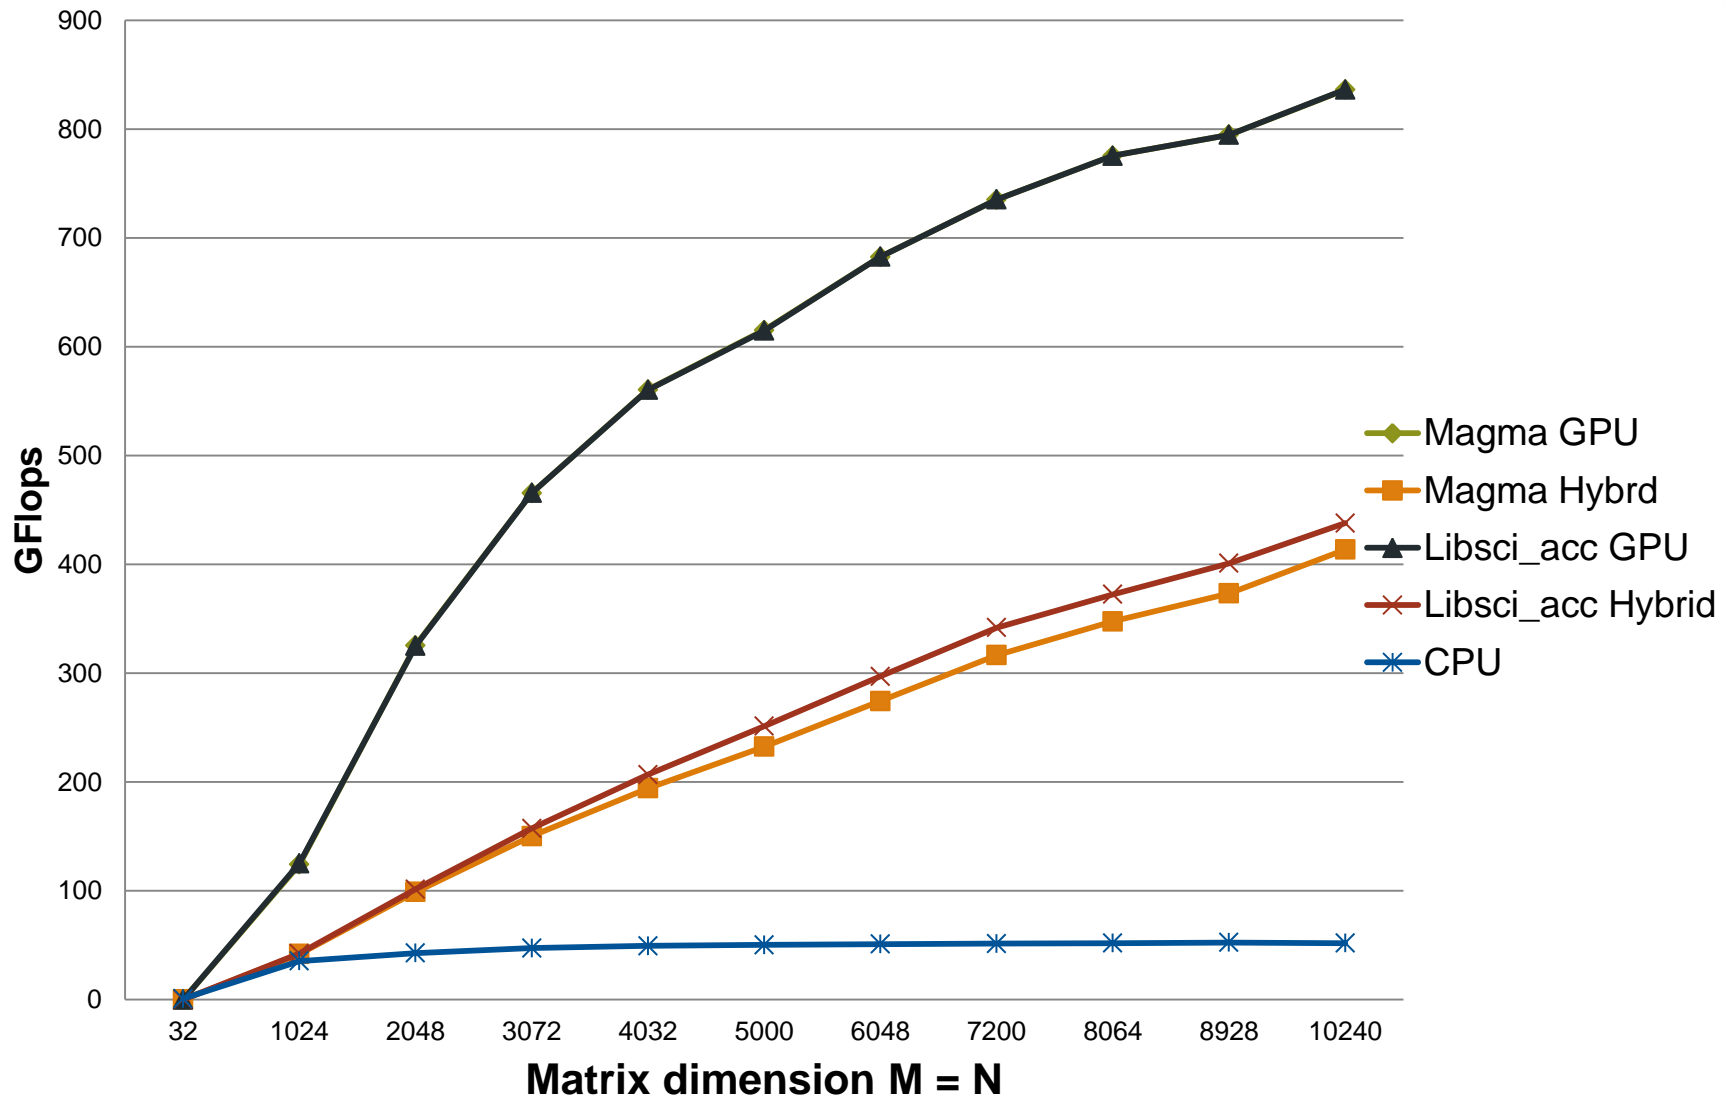
LAPACK Cholesky factorization :: double (DPOTRF)

XK7 Kepler :: Nov 2012

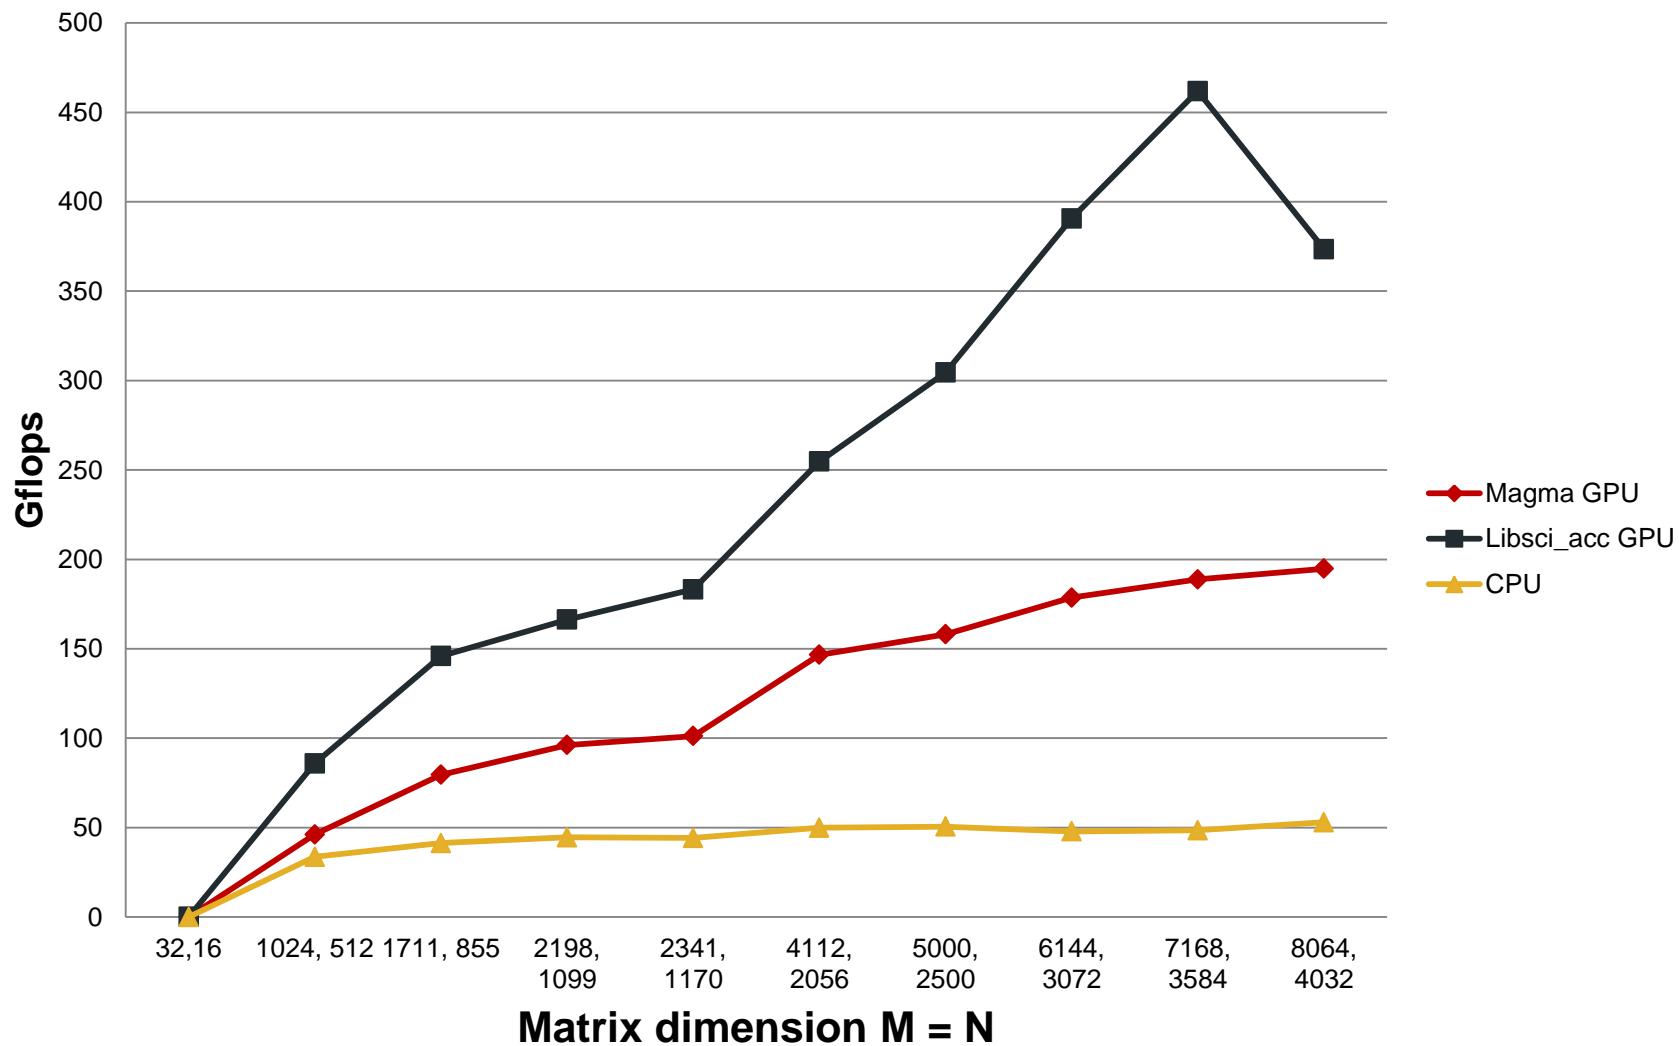


LAPACK Cholesky factorization :: double complex (ZPOTRF)

XK7 Kepler :: Nov 2012

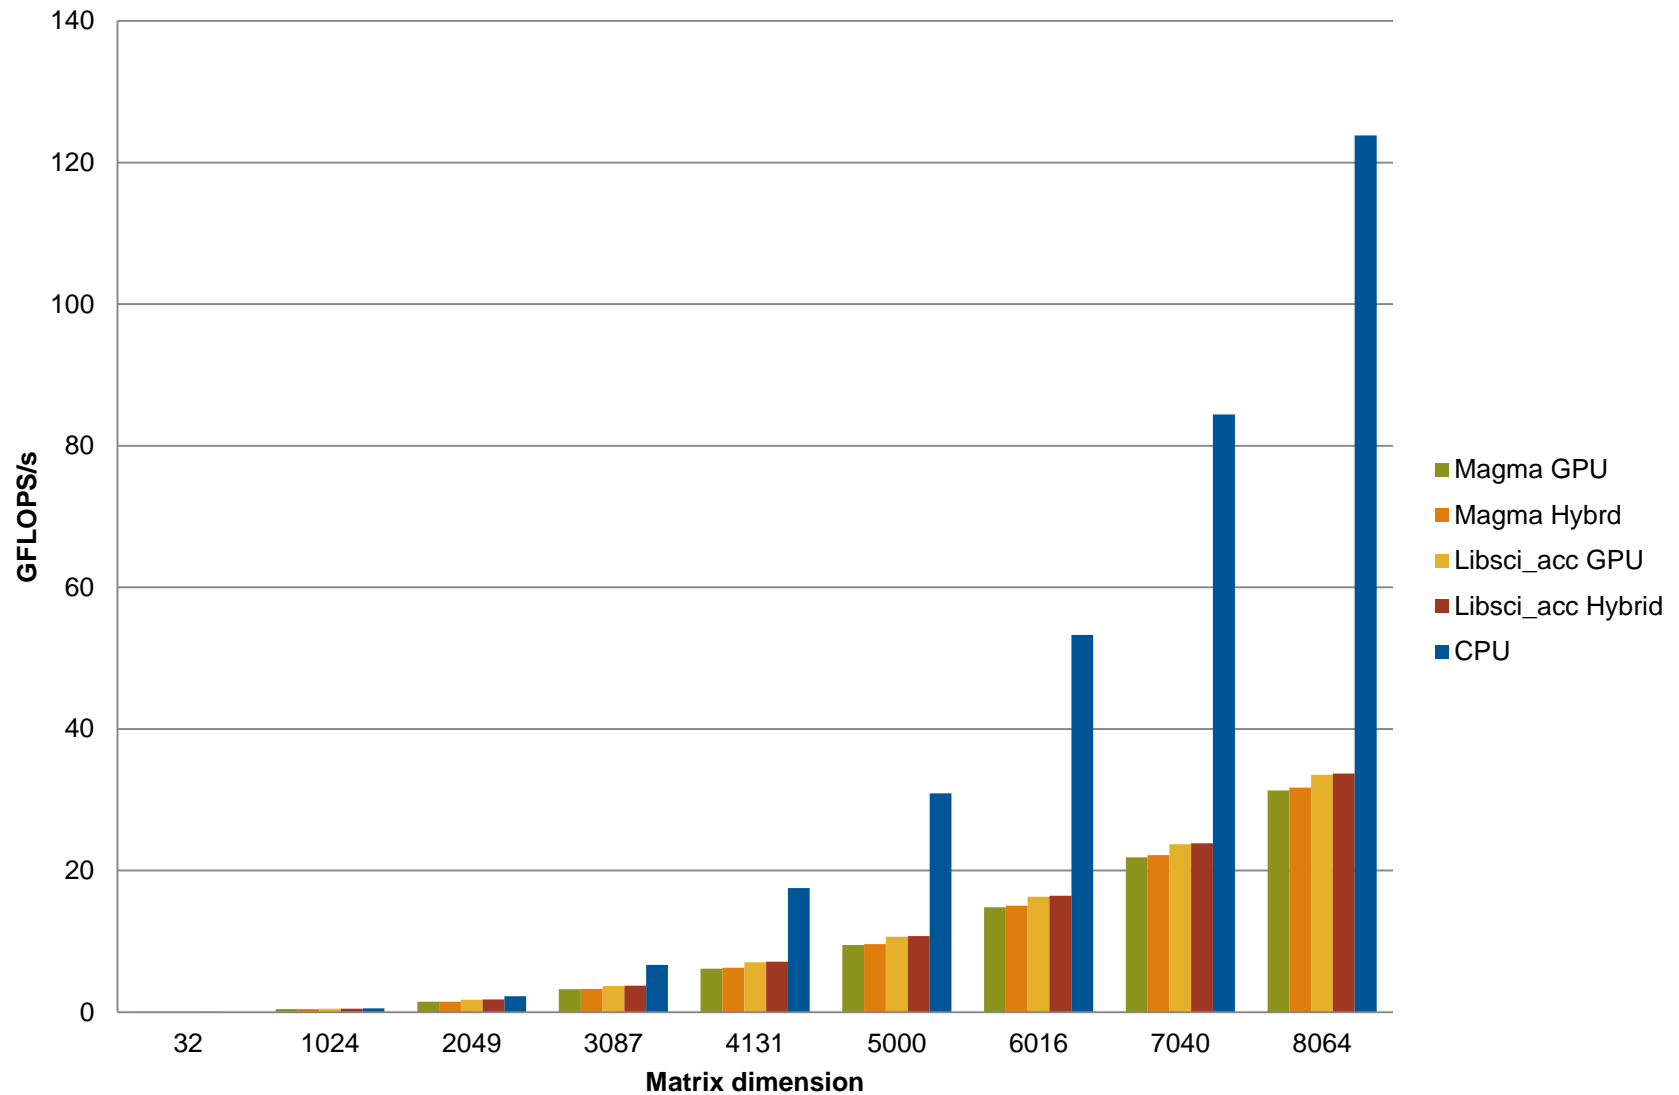


LAPACK Cholesky solver :: double (DPOTRS) XK7 Kepler :: Nov 2012



LAPACK divide and conquer eigensolver :: double (DSYEVD)

XK7 Kepler :: Nov 2012



libsci_acc BLAS Routines Available

- **BLAS 3 - HYBRID Implementations**
 - [s,d,c,z]GEMM
 - [s,d,c,z]GEMM
 - [s,d,c,z]TRSM
 - [z,c]HEMM
 - [s,d,c,z]SYMM
 - [s,d,c,z]SYRK
 - [z,d]HERK
 - [s,d,c,z]SYR2K
 - [s,d,c,z]TRMM

- **The following are supported without HYBRID implementations because there is no performance advantage**
 - All BLAS 2 Routines
 - All BLAS 1 Routines

libsci_acc LAPACK Routines Available

● HYBRID Implementations:

- [d,z]GETRF (LU Factorization)
- [d,z]POTRF (Cholesky Factorization)
- [d,z]GETRS (System Solver)
- [d,z]POTRS (System Solver)
- [d,z]GESDD* (Generalized Singular Values)
- [d,z]GEBRD (Generalized Bidiagonalization)
- [d,z]GEQRF* (QR Factorization)
- [d,z]GELQF (LQ Factorization)
- [d,z]GEEV (Non-symmetric Eigenvalues)
- DSYEVR* / ZHEEVR* (Hermitian/Symmetric Eigenvalues)
- DSYEV / DSYEVD (Hermitian/Symmetric Eigenvalues)
- ZHEEV / ZHEEVD (Hermitian/Symmetric Eigenvalues)
- DSYGVD / ZHEGVD (Hermitian/Symmetric Eigenvalue System Solver)

* Include Cray Proprietary Optimizations