



OpenACC 2.0 Highlights

Jeff Larkin, NVIDIA

OpenACC 2.0 - Highlights

- **Procedure calls, separate compilation**
- **Nested parallelism (support for Dynamic Parallelism)**
- **Loop tile clause**
- **Data management features and global data**
- **Device-specific tuning**
- **Asynchronous behavior additions**
- **New API routines**
- **New atomic construct**
- **New default(none) data clause**

Procedure calls – OpenACC 1.0

```
#pragma acc parallel loop
for ( int i=0; i<n; ++i) {
    foo(v,i);
    //must inline foo
}
```

Procedure calls – OpenACC 2.0

```
#pragma acc routine worker  
extern void foo(float* v, int i);
```

```
#pragma acc parallel loop  
for ( int i=0; i<n; ++i) {  
    foo(v,i);  
    //call on the device  
}
```



Tell the compiler the
level of parallelism
in foo

Procedure calls – OpenACC 2.0

```
#pragma acc routine worker
extern void foo(float* v,int i);

#pragma acc parallel loop
for ( int i=0; i<n; ++i) {
    foo(v,i);
    //call on the device
}
```

```
#pragma acc routine worker
void foo(float* v, int i) {
    #pragma acc loop worker
    for ( int j=0; j<n; ++j) {
        v[i*n+j] = 1.0f/(i*j);
    }
}
```

Nested Parallelism

```
#pragma acc routine
extern void foo(float* v,int i);

#pragma acc parallel loop
for ( int i=0; i<n; ++i) {
    foo(v,i);
    //call on the device and span new
    //threads
}
```

```
#pragma acc routine
void foo(float* v, int i) {
    #pragma acc parallel loop
    for ( int j=0; j<n; ++j) {
        v[i*n+j] = 1.0f/(i*j);
    }
}
```

Loop tile Clause

Block loops into
8x8 tiles.

- OpenACC 1.0 does not provide a standard way to decompose loops into 2D threadblocks
 - May better exploit data locality

```
#pragma acc loop tile(8,8)
for ( int i=0; i<n; ++i)
{
    for ( int j=0; j<n; ++j)
    {
        v[i*n+j] = 1.0f/(i*j);
    }
}
```

Data management – global data

```
float a[1000000];
```

```
extern float a[];
```

```
extern void foo(float* v,int i);
```

```
for ( int i=0; i<n; ++i) {  
    foo(v,i);  
}
```

```
void foo(float* v, int i) {
```

```
    for ( int j=0; j<n; ++j) {  
        v[i*n+j] = a[i]/(i*j);  
    }  
}
```


Data management – global data

```
float a[1000000];  
#pragma acc declare create(a)  
  
#pragma acc routine worker  
extern void foo(float* v,int i);  
  
#pragma acc parallel loop  
for ( int i=0; i<n; ++i) {  
    foo(v,i);  
    //call on the device  
}
```

```
extern float a[];  
#pragma acc declare create(a)  
  
#pragma acc routine worker  
void foo(float* v, int i) {  
    #pragma acc loop worker  
    for ( int j=0; j<n; ++j) {  
        v[i*n+j] = a[i]/(i*j);  
    }  
}
```

Data management – global data

```
float a[1000000];  
#pragma acc declare device_resident(a)  
  
#pragma acc routine worker  
extern void foo(float* v,int i);  
  
#pragma acc parallel loop  
for ( int i=0; i<n; ++i) {  
    foo(v,i);  
    //call on the device  
}
```

```
extern float a[];  
#pragma acc declare device_resident(a)  
  
#pragma acc routine worker nohost  
void foo(float* v, int i) {  
    #pragma acc loop worker  
    for ( int j=0; j<n; ++j) {  
        v[i*n+j] = a[i]/(i*j);  
    }  
}
```

Data management – structured data lifetime (OpenACC 1.0)

```
#pragma acc data copyin(x[0:n]) \  
    create(y[0:n])  
{  
}
```

Data management – unstructured data lifetime

```
class Matrix {  
    Matrix() {  
        v = new double[n];  
  
    }  
    ~Matrix() {  
  
        delete[] v;}  
private:  
    double* v;  
}
```

```
class Matrix {  
    Matrix() {  
        v = new double[n];  
        #pragma acc enter data \  
            create(v[0:n])  
    }  
    ~Matrix() {  
        #pragma acc exit data \  
            delete(v[0:n])  
        delete[] v;}  
private:  
    double* v;  
}
```

Device-specific tuning – device_type

```
#pragma acc routine worker
extern void foo(float* v,int i);
```

```
#pragma acc parallel loop \
    num_workers(384)
```

```
for ( int i=0; i<n; ++i) {
    foo(v,i);
}
```

```
#pragma acc routine worker
extern void foo(float* v,int i);
```

```
#pragma acc parallel loop \
    device_type(nvidia) num_workers(256) \
    device_type(radeon) num_workers(512)
```

```
for ( int i=0; i<n; ++i) {
    foo(v,i);
}
```

OpenACC 1.0 - **async** clause

```
#pragma acc parallel async(1)  
{... /*kernel A*/}
```

```
do_something_on_host()
```

```
#pragma acc parallel async(2)  
{.../*Kernel B*/}
```

```
#pragma acc parallel async(2)  
{.../*Kernel C*/}
```

The **async** clause is optional on the **parallel** and **kernels** constructs; when there is no **async** clause, the host process will wait until the parallel or kernels region is complete before executing any of the code that follows the construct. When there is an **async** clause,

Do not wait for kernel completion

Executes concurrently with kernels

(potentially) executes concurrently with kernels

OpenACC 1.0 – **wait** directive

```
#pragma acc parallel async(1)
{... /*kernel A*/}
```

```
do_something_on_host()
```

```
#pragma acc parallel async(2)
{.../*Kernel B*/}
```

```
#pragma acc parallel async(2)
{.../*Kernel C*/}
```

```
#pragma acc wait(1)
#pragma acc parallel async(2)
{...}
```

Wait for kernel A in queue (stream) 1 - blocks host

Schedule new work in queue (stream) 2 that depends on queue 1.

OpenACC 2.0 – wait directive with async clause

```
#pragma acc parallel async(1)
{... /*kernel A*/}
```

```
do_something_on_host()
```

```
#pragma acc parallel async(2)
{.../*Kernel B*/}
```

```
#pragma acc parallel async(2)
{.../*Kernel C*/}
```

```
#pragma acc wait(1)
#pragma acc parallel async(2)
{...}
```

```
#pragma acc parallel async(1)
{... /*kernel A*/}
```

```
do_something_on_host()
```

```
#pragma acc parallel async(2)
{.../*Kernel B*/}
```

```
#pragma acc parallel async(2)
{.../*Kernel C*/}
```

```
#pragma acc wait(1) async(2)
#pragma acc parallel async(2)
{...}
```


New API routines

- **Improved Data management for C/C++**
 - `acc_copyin`
 - `acc_present_or_copyin`
 - `acc_create`
 - `acc_present_or_create`
 - `acc_copyout`
 - `acc_delete`
 - `acc_map_data`
 - `acc_unmap_data`
 - `acc_deviceptr`
 - `acc_hostptr`
 - `acc_is_present`
 - `acc_memcpy_to_device`
 - `acc_memcpy_from_device`
 - `acc_update_device`
 - `acc_update_self`
- **Expanded Interoperability with CUDA, OpenCL, & Xeon Phi**
 - `acc_get_cuda_stream`
 - `acc_get_current_cuda_device`
 - ...

OpenACC 2.0 - Highlights

- **Procedure calls, separate compilation**
- **Nested parallelism (support for Dynamic Parallelism)**
- **Loop tile clause**
- **Data management features and global data**
- **Device-specific tuning**
- **Asynchronous behavior additions**
- **New API routines**
- **New atomic construct**
- **New default(none) data clause**