


Performance Analysis with Vampir

17 July 2013


Joseph Schuchart

Vampir On-Site Support



The Vampir Toolsuite

- Scalable Trace-based performance analysis tools
- Detailed insight into parallel applications
- Two components:
 - **VampirTrace**: Instrumentation and Measurement
 - **Vampir**: Visual Analysis of trace data



VampirTrace

- **Instrumentation**

- Insert **function-level instrumentation** (Compiler, TAU, Dyninst, Manual)
- Provide **instrumented libraries**:
 - CUDA / CUPTI
 - MPI
 - libc
- Instrument OpenMP (Opari)

- **Measurement:**

- Record generated events
- Manage event buffers
- Take care of post-processing
- Flexible configurable through environment variables
- Profiling and tracing

VampirTrace: Workflow

1. Instrument your code with VampirTrace

1. Load the VampirTrace module

```
module load vampirtrace  
OR  
module load vampirtrace/5.14.3-chester
```

2. Use compiler wrappers:

```
CC      = cc  
CXX     = CC  
F90     = ftn  
NVCC    = nvcc
```



```
CC      = vtcc  
CXX     = vtcxx  
F90     = vtf90  
NVCC    = vtnvcc
```

3. Recompile/Relink

2. Run your code with an appropriate test set


3. Analyze the data using Vampir

VampirTrace: Environment Variables

Environment Variable	Description	Default
VT_BUFFER_SIZE	Size of event buffer	32M
VT_MAX_FLUSHES	Maximum number of flushes	1
VT_PFORM_GDIR	Directory to hold trace data	. /
VT_METRICS	Colon-separated list of PAPI metrics to record	-
VT_FILTER_SPEC	Define a filter file for functions	-
VT_IOTRACE	Enable I/O tracing	no
VT_MEMTRACE	Enable memory tracing	no
VT_MODE	Colon-separated list of VampirTrace modes: Tracing (TRACE) or Profiling (STAT)	TRACE
VT_VERBOSE	Controls verbosity: Quiet (0), Warnings (1), Information (2)	0

VampirTrace: GPU Support

- VampirTrace supports GPU applications
 - Uses **NVIDIA CUPTI: CUDA Performance Tools Interface**
 - Low overhead for the GPU parts
 - Records GPU activities
 - Data movement (Host <-> Device)
 - Kernel runtimes
 - CUDA function calls
 - Can be used for CUDA and OpenACC applications
 - Waiting for the OpenACC tools API
 - Has to be explicitly enabled using VT_GPUTRACE



VampirTrace: VT_GPUTRACE

Option	Description
cuda	Enable support for CUDA
runtime	CUDA runtime API tracing
driver	CUDA driver API tracing
kernel	Record kernel activity
idle	Record GPU idle times
memcpy	Record GPU memory copies
stream_reuse	force reusing of CUDA streams after <code>cudaStreamDestroy()</code>
memusage	Record GPU memory usage
yes default	Same as “cuda, runtime, kernel, memcpy”
no	Disable CUDA measurement

VampirTrace: Additional Variables

Variable	Description
VT_GPUTRACE_KERNEL=[yes 2]	yes: record kernel activity (same as kernel) 2: record additional kernel counter
VT_CUDATRACE_SYNC=[0 1 2 3]	Controls how VampirTrace handles synchronization (CUDA runtime wrapper only)
VT_CUPTI_METRICS	Record additional CUPTI metrics, e.g., VT_CUPTI_METRICS=local_store:local_load
VT_CUPTI_SAMPLING=[yes no]	Poll for CUPTI counter values during kernel execution
VT_GPUTRACE_MEMUSAGE=[yes 2]	yes: Add counter gpu_mem_usage (same as memusage) 2: Also print missing cudaFree() calls to stderr

- For CUPTI counter overview:
 - http://docs.nvidia.com/cuda/cupti/index.html#r_metric_api

VampirTrace: OpenACC

- Most compilers use the CUDA driver API:
 - Set
`VT_GPUTRACE=cuda,cupti,runtime,driver,kernel,memcpy,memoryusage,idle`
- The default CUPTI buffer size is too small (64k)
 - Set `VT_CUDATRACE_BUFFER_SIZE=64M`

Vampir: Example

1. Load Vampir module:

```
module load vampir  
OR  
module load vampir/8.1.0-130626-chester
```

2. Launch VampirServer

```
vampirserver start -a <projectID> -n <N>
```

3. Setup SSH tunnel as given by Vampir

```
ssh -L 30001:nid00809:30066 \  
<user>@titan-internal.ccs.ornl.gov
```

4. Connect your local Vampir installation

Note: X-Forwarding possible but slow!

Vampir: Hints

- Number of server processes:
 - Usually memory bound
 - Depends on trace size
 - About 4x deflate
- Swap PrgEnv before loading the Vampir module
 - PGI-built module has problems with MPI
- Consider using Lens for analysis

How to get Vampir

- ORNL has a site-wide license
 - You can download and install Vampir on your local machine
 - Activate it at www.vampir.eu/activation
- Remote-only version available if you're not on-site
- See:
https://www.olcf.ornl.gov/kb_articles/software-vampir/?softwaretype=kb_software_debugging_and_profiling

```
$ scp <user>@home.ccs.ornl.gov:/sw/sources/vampir/client/vampir-*.dmg .  
$ scp <user>@home.ccs.ornl.gov:/sw/sources/vampir/client/vampir.license .  
# Proceed with dmg installation by opening it and dragging the Vampir  
  icon to your preferred launch spot. Start Vampir using this icon.
```

VampirTrace: Hints

- Traces can become huge:
 - Esp. C++ problematic
 - Solution: Filter events
 - Runtime filter: reduce trace size but do not reduce overhead
 - TAU: Compile time filter
 - GCC (just in case ;)) instrumentation filter

```
1 MPI_* -- -1
2 std::* -- 0
3 * -- 1000
```

```
BEGIN_EXCLUDE_LIST
init_*
boost::*
END_EXCLUDE_LIST

BEGIN_FILE_EXCLUDE_LIST
/usr/include/*
END_FILE_EXCLUDE_LIST
```

```
-finstrument-functions-exclude-file-list=include/g++,math_extra
-finstrument-functions-exclude-function-list=map,timing,Timer,operator
```

VampirTrace: Hints

- Use profiles to determine heavily called functions
 - Set `VT_MODE=STAT`
- The default buffer size is quite small:
 - Consider setting `VT_BUFFER_SIZE=512M`
 - Allocates 10% of that per thread
 - Override with `VT_THREAD_BUFFER_SIZE`
 - Also true for CUPTI events:
 - Override with `VT_CUDATRACE_BUFFER_SIZE`
- The compiler wrapper doesn't recognize the parallelization paradigm:
 - Use `-vt:mpi`, `-vt:mt`, or `-vt:hyb`

VampirTrace: Hints

- Take caution with large core counts:
 - Produces a large number of files (two per process/thread)
 - Can cause problems with Lustre
- VT comes with an I/O forwarding solution
 - Allocate additional nodes for I/O server
 - Launch server in your job script
 - Trace data will be written through the server
 - Only two files per server

```
module load vampirtrace
source vtiofsl-start.sh -n M
aprun -n N a.out
source vtiofsl-stop.sh
```

Vampir: Live demo